

Advanced Databases

Query Optimization

Foundations and Trends in Databases

J. M. Hellerstein, M. Stonebraker
and J. Hamilton

Vol. 1, No. 2 (2007) 141–259, 2007


DOI: 10.1561/19000000002



Query Optimizer


- Three components:
 - 1. Search space
 - 2. Plan enumeration algorithms
 - 3. Cardinality and cost estimation

- Note: First query optimizer was for System R, from IBM, in 1979, you will often see this references when reading about optimization in relational databases




Approaches

- Heuristics
- Heuristics + Cost-based Join Order Search
- Randomized Algorithms
- Stratified Search
- Unified Search



Heuristic Optimization

- Define static rules that transform logical operators to a physical plan.
 - Perform most restrictive selection early
 - Perform all selections before joins
 - Predicate/Limit/Projection pushdowns
 - Predicate part of SQL that filters data
 - Pushdown
 - If you issue a query in one place to run against a lot of data that's in another place, you could spawn a lot of network traffic, which could be slow and costly.
 - If you can “push down” parts of the query to where the data is stored, and thus filter out most of the data, then you can greatly reduce network traffic.
- Join ordering based on cardinality



Example

```
CREATE TABLE ARTIST (  
  ID INT PRIMARY KEY,  
  NAME VARCHAR(32)  
);
```

```
CREATE TABLE ALBUM (  
  ID INT PRIMARY KEY,  
  NAME VARCHAR(32) UNIQUE  
);
```

```
CREATE TABLE APPEARS (  
  ARTIST_ID INT  
    ↪REFERENCES ARTIST(ID),  
  ALBUM_ID INT  
    ↪REFERENCES ALBUM(ID),  
  PRIMARY KEY  
    ↪(ARTIST_ID, ALBUM_ID)  
);
```

Example


Retrieve the names of people that appear on Joy's mixtape

```
SELECT ARTIST.NAME  
FROM ARTIST, APPEARS, ALBUM  
WHERE ARTIST.ID=APPEARS.ARTIST_ID  
AND APPEARS.ALBUM_ID=ALBUM.ID  
AND ALBUM.NAME="Joy's Covid Remix"
```

Example

Retrieve the names of people that appear on Joy's mixtape

```
SELECT ARTIST.NAME  
  FROM ARTIST, APPEARS, ALBUM  
 WHERE ARTIST.ID=APPEARS.ARTIST_ID  
       AND APPEARS.ALBUM_ID=ALBUM.ID  
       AND ALBUM.NAME="Joy's Covid Remix"
```



*Step #1: Decompose into
single-variable queries*

Q1

```
SELECT ALBUM.ID AS ALBUM_ID INTO TEMP1  
  FROM ALBUM  
 WHERE ALBUM.NAME="Joy's Covid Remix"
```


Q2

```
SELECT ARTIST.NAME  
  FROM ARTIST, APPEARS, TEMP1  
 WHERE ARTIST.ID=APPEARS.ARTIST_ID  
       AND APPEARS.ALBUM_ID=TEMP1.ALBUM_ID
```

Example

Retrieve the names of people that appear on Joy's mixtape

```
SELECT ARTIST.NAME  
  FROM ARTIST, APPEARS, ALBUM  
 WHERE ARTIST.ID=APPEARS.ARTIST_ID  
       AND APPEARS.ALBUM_ID=ALBUM.ID  
       AND ALBUM.NAME="Joy's Covid Remix"
```



*Step #1: Decompose into
single-variable queries*

Q1

```
SELECT ALBUM.ID AS ALBUM_ID INTO TEMP1  
  FROM ALBUM  
 WHERE ALBUM.NAME="Joy's Covid Remix"
```

Q3

```
SELECT APPEARS.ARTIST_ID INTO TEMP2  
  FROM APPEARS, TEMP1  
 WHERE APPEARS.ALBUM_ID=TEMP1.ALBUM_ID
```


Q4

```
SELECT ARTIST.NAME  
  FROM ARTIST, TEMP2  
 WHERE ARTIST.ARTIST_ID=TEMP2.ARTIST_ID
```


Example

Retrieve the names of people that appear on Joy's mixtape

```
SELECT ARTIST.NAME  
  FROM ARTIST, APPEARS, ALBUM  
 WHERE ARTIST.ID=APPEARS.ARTIST_ID  
       AND APPEARS.ALBUM_ID=ALBUM.ID  
       AND ALBUM.NAME="Joy's Covid Remix"
```



*Step #1: Decompose into
single-variable queries*

*Step #2: Substitute the values
from Q1→Q3→Q4*

Q1

```
SELECT ALBUM.ID AS ALBUM_ID INTO TEMP1  
  FROM ALBUM  
 WHERE ALBUM.NAME="Joy's Covid Remix"
```

Q3

```
SELECT APPEARS.ARTIST_ID INTO TEMP2  
  FROM APPEARS, TEMP1  
 WHERE APPEARS.ALBUM_ID=TEMP1.ALBUM_ID
```


Q4

```
SELECT ARTIST.NAME  
  FROM ARTIST, TEMP2  
 WHERE ARTIST.ARTIST_ID=TEMP2.ARTIST_ID
```

Example

Retrieve the names of people that appear on Joy's mixtape

```
SELECT ARTIST.NAME  
  FROM ARTIST, APPEARS, ALBUM  
 WHERE ARTIST.ID=APPEARS.ARTIST_ID  
       AND APPEARS.ALBUM_ID=ALBUM.ID  
       AND ALBUM.NAME="Joy's Covid Remix"
```



*Step #1: Decompose into
single-variable queries*

*Step #2: Substitute the values
from Q1→Q3→Q4*

Q1

```
SELECT ALBUM.ID AS ALBUM_ID INTO TEMP1  
  FROM ALBUM  
 WHERE ALBUM.NAME="Joy's Covid Remix"
```

Q3

```
SELECT APPEARS.ARTIST_ID INTO TEMP2  
  FROM APPEARS, TEMP1  
 WHERE APPEARS.ALBUM_ID=TEMP1.ALBUM_ID
```

Q4

```
SELECT ARTIST.NAME  
  FROM ARTIST, TEMP2  
 WHERE ARTIST.ARTIST_ID=TEMP2.ARTIST_ID
```

Example

Retrieve the names of people that appear on Joy's mixtape

```
SELECT ARTIST.NAME  
  FROM ARTIST, APPEARS, ALBUM  
 WHERE ARTIST.ID=APPEARS.ARTIST_ID  
       AND APPEARS.ALBUM_ID=ALBUM.ID  
       AND ALBUM.NAME="Joy's Covid Remix"
```



ALBUM_ID
9999

*Step #1: Decompose into
single-variable queries*

*Step #2: Substitute the values
from Q1→Q3→Q4*

```
SELECT APPEARS.ARTIST_ID  
  FROM APPEARS  
 WHERE APPEARS.ALBUM_ID=9999
```

Q4

```
SELECT ARTIST.NAME  
  FROM ARTIST, TEMP2  
 WHERE ARTIST.ARTIST_ID=TEMP2.ARTIST_ID
```

Example

Retrieve the names of people that appear on Joy's mixtape

```
SELECT ARTIST.NAME  
  FROM ARTIST, APPEARS, ALBUM  
 WHERE ARTIST.ID=APPEARS.ARTIST_ID  
       AND APPEARS.ALBUM_ID=ALBUM.ID  
       AND ALBUM.NAME="Joy's Covid Remix"
```



ALBUM_ID
9999

ARTIST_ID
123
456

*Step #1: Decompose into
single-variable queries*

*Step #2: Substitute the values
from Q1→Q3→Q4*

Q4

```
SELECT ARTIST.NAME  
  FROM ARTIST, TEMP2  
 WHERE ARTIST.ARTIST_ID=TEMP2.ARTIST_ID
```

Example

Retrieve the names of people that appear on Joy's mixtape

```
SELECT ARTIST.NAME  
  FROM ARTIST, APPEARS, ALBUM  
 WHERE ARTIST.ID=APPEARS.ARTIST_ID  
       AND APPEARS.ALBUM_ID=ALBUM.ID  
       AND ALBUM.NAME="Joy's Covid Remix"
```



*Step #1: Decompose into
single-variable queries*

*Step #2: Substitute the values
from Q1→Q3→Q4*

ALBUM_ID
9999

ARTIST_ID
123
456

```
SELECT ARTIST.NAME  
  FROM ARTIST  
 WHERE ARTIST.ARTIST_ID=123
```

```
SELECT ARTIST.NAME  
  FROM ARTIST  
 WHERE ARTIST.ARTIST_ID=456
```

Example

Retrieve the names of people that appear on Joy's mixtape

```
SELECT ARTIST.NAME  
FROM ARTIST, APPEARS, ALBUM  
WHERE ARTIST.ID=APPEARS.ARTIST_ID  
AND APPEARS.ALBUM_ID=ALBUM.ID  
AND ALBUM.NAME="Joy's Covid Remix"
```



*Step #1: Decompose into
single-variable queries*

*Step #2: Substitute the values
from Q1→Q3→Q4*

ALBUM_ID
9999


ARTIST_ID
123
456

NAME
O.D.B.

NAME
DJ Premier


Heuristic Optimization

- Advantages:
 - Easy to implement and debug.
 - Works reasonably well and is fast for simple queries.
- Disadvantages:
 - Relies on magic constants that predict the efficacy of a planning decision.
 - Magic constants e.g. block size, parallelism granularity, thread weight
 - Nearly impossible to generate good plans when operators have complex inter-dependencies.




Heuristic Optimization + Cost Estimation

- Use static rules to perform initial optimization.
- Then use dynamic programming to determine the best join order for tables.
 - First cost-based query optimizer
 - Bottom-up planning (forward chaining) using a divide-and-conquer search method



System R Optimizer

- Break query up into blocks and generate the logical operators for each block.
- For each logical operator, generate a set of physical operators that implement it.
 - All combinations of join algorithms and access paths
- Then iteratively construct a “left-deep” tree that minimizes the estimated amount of work to execute the plan.



Example

Retrieve the names of people that appear on Joy's mixtape ordered by their artist id.

```
SELECT ARTIST.NAME  
FROM ARTIST, APPEARS, ALBUM  
WHERE ARTIST.ID=APPEARS.ARTIST_ID  
AND APPEARS.ALBUM_ID=ALBUM.ID  
AND ALBUM.NAME="Joy's Covid Remix"  
ORDER BY ARTIST.ID
```

Example

Retrieve the names of people that appear on Joy's mixtape ordered by their artist id.

```
SELECT ARTIST.NAME
FROM ARTIST, APPEARS, ALBUM
WHERE ARTIST.ID=APPEARS.ARTIST_ID
AND APPEARS.ALBUM_ID=ALBUM.ID
AND ALBUM.NAME="Joy's Covid Remix"
ORDER BY ARTIST.ID
```

Step #1: Choose the best access paths to each table

Step #2: Enumerate all possible join orderings for tables

Step #3: Determine the join ordering with the lowest cost

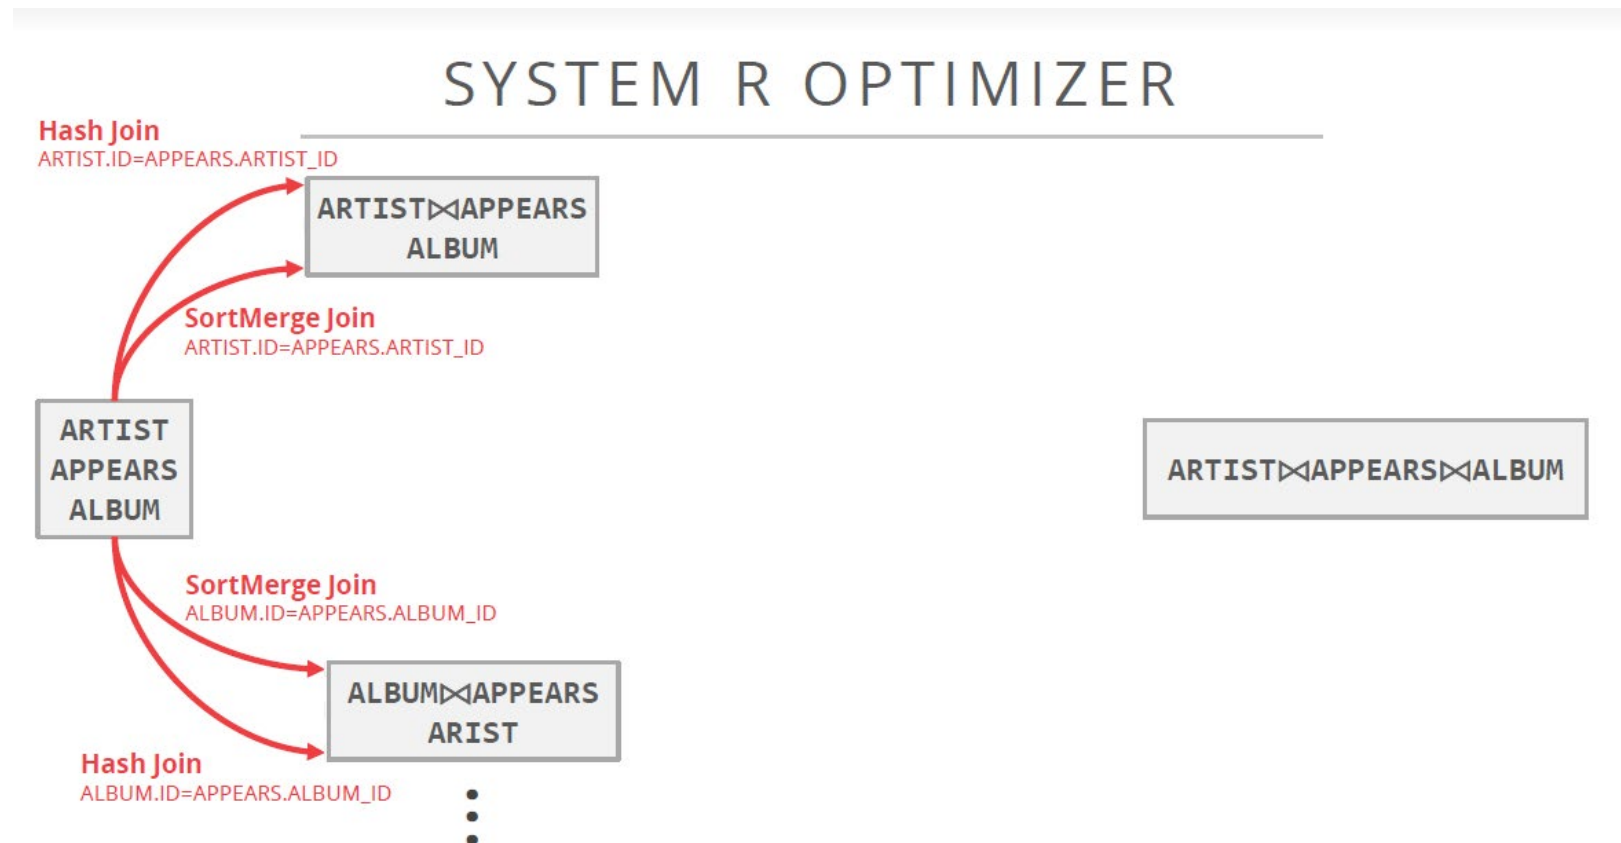
ARTIST: Sequential Scan

APPEARS: Sequential Scan

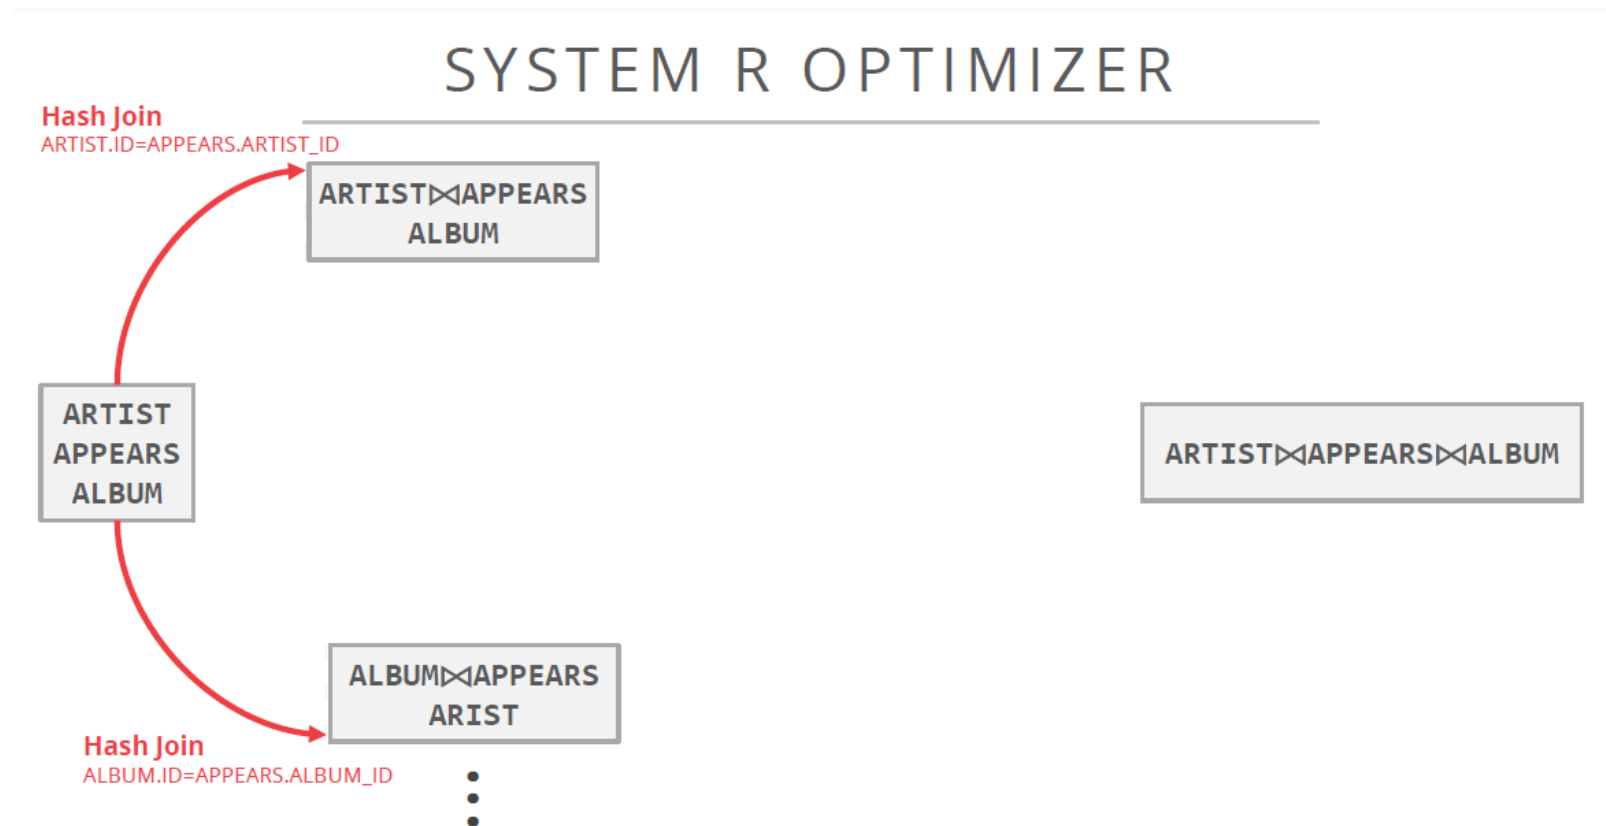
ALBUM: Index Look-up on NAME

ARTIST	⋈	APPEARS	⋈	ALBUM
APPEARS	⋈	ALBUM	⋈	ARTIST
ALBUM	⋈	APPEARS	⋈	ARTIST
APPEARS	⋈	ARTIST	⋈	ALBUM
ARTIST		ALBUM	⋈	APPEARS
ALBUM		ARTIST	⋈	APPEARS
⋮		⋮		⋮

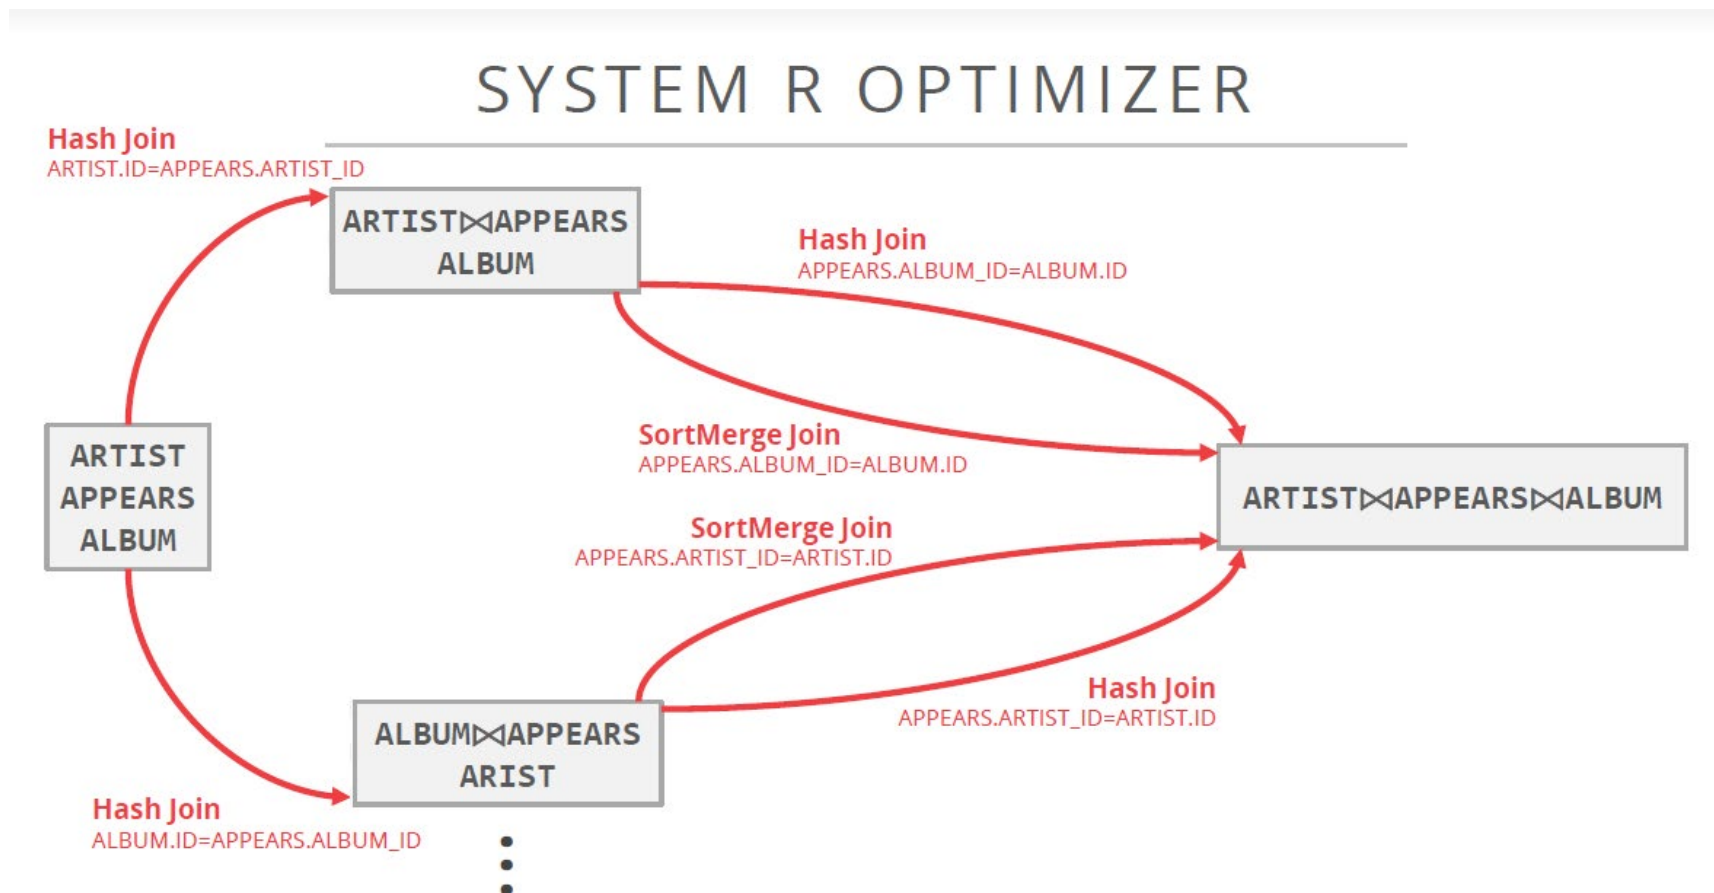
Example



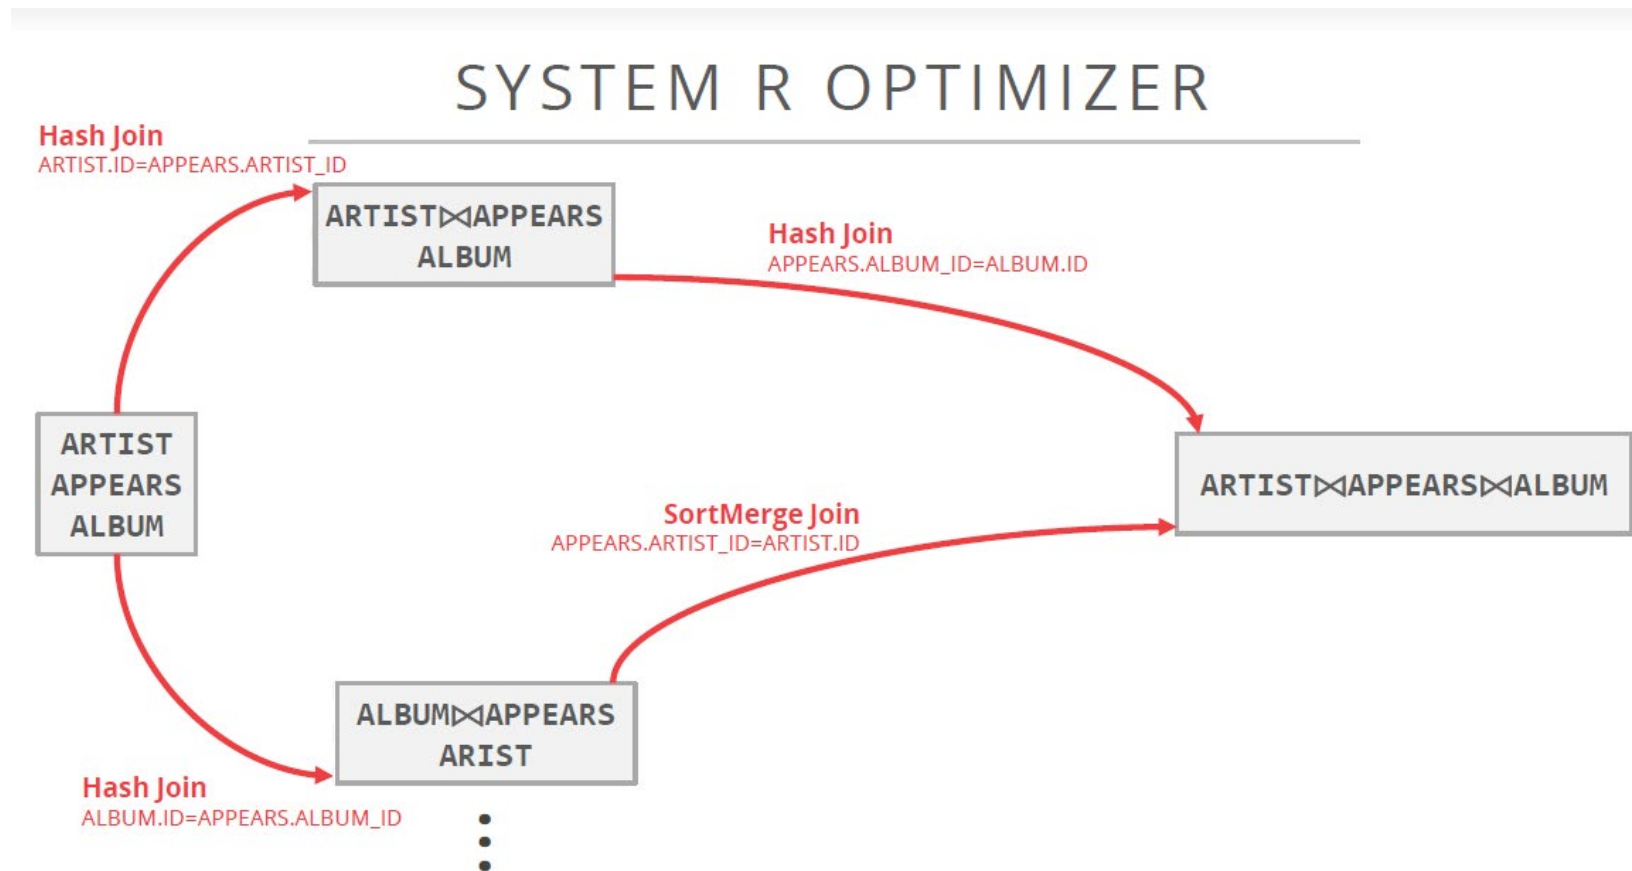
Example



Example




Example




Heuristic Optimization + Cost Estimation

- Advantages:
 - Easy to implement and debug.
 - Works reasonably well and is fast for simple queries.
- Disadvantages:
 - Relies on magic constants that predict the efficacy of a planning decision.
 - Nearly impossible to generate good plans when operators have complex inter-dependencies



Randomized Algorithms


- Perform a random walk over a solution space of all possible (valid) plans for a query
- Continue searching until a cost threshold is reached or the optimizer runs for a particular length of time
- Example: Postgres' genetic algorithm



Randomized Algorithms


Simulated Annealing

- Start with a query plan that is generated using the heuristic-only approach
- Compute random permutations of operators (e.g., swap the join order of two tables)
 - Always accept a change that reduces cost
 - Only accept a change that increases cost with some probability
 - Reject any change that violates correctness (e.g., sort ordering)




Randomized Algorithm

- Advantages:
 - Jumping around the search space randomly allows the optimizer to get out of local minimums.
 - Low memory overhead (if no history is kept).
- Disadvantages:
 - Difficult to determine why the DBMS may have chose a particular plan.
 - Have to do extra work to ensure that query plans are deterministic.
 - Still have to implement correctness rules.




Optimizer Generators

- Use a rule engine that allows transformations to modify the query plan operators.
- The physical properties of data is embedded with the operators themselves.
- Choice #1: Stratified Search
 - Planning is done in multiple stages
- Choice #2: Unified Search
 - Perform query planning all at once.




Stratified Search

- First rewrite the logical query plan using transformation rules.
 - The engine checks whether the transformation is allowed before it can be applied.
 - Cost is never considered in this step.
- Then perform a cost-based search to map the logical plan to a physical plan




Stratified Search

- Example: Starburst Optimizer
 - Better implementation of the System R optimizer that uses declarative rules.
 - Stage #1: Query Rewrite
 - Compute a SQL-block-level, relational calculus-like representation of queries.
 - Stage #2: Plan Optimization
 - Execute a System R-style dynamic programming phase once query rewrite has completed.
- Used by latest version of IBM DB2




Stratified Search

- Advantages:
 - Works well in practice with fast performance.
- Disadvantages:
 - Difficult to assign priorities to transformations
 - Some transformations are difficult to assess without computing multiple cost estimations.
 - Rules maintenance is a huge pain.




Unified Search

- Unify the notion of both logical-logical and logical-physical transformations.
 - No need for separate stages because everything is transformations.
- This approach generates a lot more transformations so it makes heavy use of memorization to reduce redundant work



Unified Search

- Example: Volcano Optimizer
 - General purpose cost-based query optimizer, based on equivalence rules on algebras.
 - Easily add new operations and equivalence rules.
 - Treats physical properties of data as first-class entities during planning.
 - Top-down approach (backward chaining) using branch-and-bound search.
- Used by NonStop SQL



Unified Search

- Example: Cascades Optimizer
 - Object-oriented implementation of the Volcano query optimizer.
 - Simplistic expression re-writing can be through a direct mapping function rather than an exhaustive search.
- Used by SQL Server, Greenplum's Orca, and Apache Calcite.

